

▶ [Developer guides](#) / Customizing what happens in `fit()` with TensorFlow

Customizing what happens in `fit()` with TensorFlow

Author: [fchollet](#)

Date created: 2020/04/15

Last modified: 2023/06/27

Description: Overriding the training step of the Model class with TensorFlow.

» [View in Colab](#) · ⌂ [GitHub source](#)

[Customizing what happens in `fit\(\)` with TensorFlow](#)

[Introduction](#)

[Setup](#)

[A first simple example](#)

[Going lower-level](#)

[Supporting `sample_weight` & `class_weight`](#)

[Providing your own evaluation step](#)

[Wrapping up: an end-to-end GAN example](#)

Introduction

When you're doing supervised learning, you can use `fit()` and everything works smoothly.

When you need to take control of every little detail, you can write your own training loop entirely from scratch.

But what if you need a custom training algorithm, but you still want to benefit from the convenient features of `fit()`, such as callbacks, built-in distribution support, or step fusing?

A core principle of Keras is **progressive disclosure of complexity**. You should always be able to get into lower-level workflows in a gradual way. You shouldn't fall off a cliff if the high-level functionality doesn't exactly match your use case. You should be able to gain more control over the small details while retaining a commensurate amount of high-level convenience.

When you need to customize what `fit()` does, you should **override the training step function of the Model class**. This is the function that is called by `fit()` for every batch of data. You will then be able to call `fit()` as usual – and it will be running your own learning algorithm.

Note that this pattern does not prevent you from building models with the Functional API. You can do this whether you're building `Sequential` models, Functional API models, or subclassed models.

Let's see how that works.

Setup

```
import os

# This guide can only be run with the TF backend.
os.environ["KERAS_BACKEND"] = "tensorflow"

import tensorflow as tf
import keras
from keras import layers
import numpy as np
```

A first simple example

Let's start from a simple example:

- We create a new class that subclasses `keras.Model`.



The input argument `data` is what gets passed to fit as training data.

- If you pass NumPy arrays, by calling `fit(x, y, ...)`, then `data` will be the tuple `(x, y)`
- If you pass a `tf.data.Dataset`, by calling `fit(dataset, ...)`, then `data` will be what gets yielded by `dataset` at each batch.

In the body of the `train_step()` method, we implement a regular training update, similar to what you are already familiar with. Importantly, we compute the loss via `self.compute_loss()`, which wraps the loss(es) function(s) that were passed to `compile()`.

Similarly, we call `metric.update_state(y, y_pred)` on metrics from `self.metrics`, to update the state of the metrics that were passed in `compile()`, and we query results from `self.metrics` at the end to retrieve their current value.

```
class CustomModel(keras.Model):
    def train_step(self, data):
        # Unpack the data. Its structure depends on your model and
        # on what you pass to `fit()`.

        x, y = data

        with tf.GradientTape() as tape:
            y_pred = self(x, training=True) # Forward pass
            # Compute the loss value
            # (the loss function is configured in `compile()`)
            loss = self.compute_loss(y=y, y_pred=y_pred)

            # Compute gradients
            trainable_vars = self.trainable_variables
            gradients = tape.gradient(loss, trainable_vars)

            # Update weights
            self.optimizer.apply(gradients, trainable_vars)

            # Update metrics (includes the metric that tracks the loss)
            for metric in self.metrics:
                if metric.name == "loss":
                    metric.update_state(loss)
                else:
                    metric.update_state(y, y_pred)

        # Return a dict mapping metric names to current value
        return {m.name: m.result() for m in self.metrics}
```

Let's try this out:

```
# Construct and compile an instance of CustomModel
inputs = keras.Input(shape=(32,))
outputs = keras.layers.Dense(1)(inputs)
model = CustomModel(inputs, outputs)
model.compile(optimizer="adam", loss="mse", metrics=["mae"])

# Just use `fit` as usual
x = np.random.random((1000, 32))
y = np.random.random((1000, 1))
model.fit(x, y, epochs=3)
```

```
0.3776
Epoch 2/3
32/32 ━━━━━━━━━━━━━━━━ 0s 318us/step - mae: 0.3986 - loss:
0.2466
Epoch 3/3
32/32 ━━━━━━━━━━━━━━━━ 0s 372us/step - mae: 0.3848 - loss:
0.2319

WARNING: All log messages before absl::InitializeLog() is called are written
to STDERR
I0000 00:00:1699222602.443035      1 device_compiler.h:187] Compiled cluster
using XLA! This line is logged at most once for the lifetime of the process.

<keras.src.callbacks.history.History at 0x2a5599f00>
```

Going lower-level

Naturally, you could just skip passing a loss function in `compile()`, and instead do everything *manually* in `train_step`. Likewise for metrics.

Here's a lower-level example, that only uses `compile()` to configure the optimizer:

- We start by creating `Metric` instances to track our loss and a MAE score (in `__init__()`).
- We implement a custom `train_step()` that updates the state of these metrics (by calling `update_state()` on them), then query them (via `result()`) to return their current average value, to be displayed by the progress bar and to be pass to any callback.
- Note that we would need to call `reset_states()` on our metrics between each epoch!
Otherwise calling `result()` would return an average since the start of training, whereas we usually work with per-epoch averages. Thankfully, the framework can do that for us: just list any metric you want to reset in the `metrics` property of the model. The model will call `reset_states()` on any object listed here at the beginning of each `fit()` epoch or at the beginning of a call to `evaluate()`.



```
super().__init__(**args, **kwargs)
self.loss_tracker = keras.metrics.Mean(name="loss")
self.mae_metric = keras.metrics.MeanAbsoluteError(name="mae")
self.loss_fn = keras.losses.MeanSquaredError()

def train_step(self, data):
    x, y = data

    with tf.GradientTape() as tape:
        y_pred = self(x, training=True) # Forward pass
        # Compute our own loss
        loss = self.loss_fn(y, y_pred)

    # Compute gradients
    trainable_vars = self.trainable_variables
    gradients = tape.gradient(loss, trainable_vars)

    # Update weights
    self.optimizer.apply(gradients, trainable_vars)

    # Compute our own metrics
    self.loss_tracker.update_state(loss)
    self.mae_metric.update_state(y, y_pred)
    return {
        "loss": self.loss_tracker.result(),
        "mae": self.mae_metric.result(),
    }

@property
def metrics(self):
    # We list our `Metric` objects here so that `reset_states()` can be
    # called automatically at the start of each epoch
    # or at the start of `evaluate()`.

    return [self.loss_tracker, self.mae_metric]

# Construct an instance of CustomModel
inputs = keras.Input(shape=(32,))
outputs = keras.layers.Dense(1)(inputs)
model = CustomModel(inputs, outputs)

# We don't pass a loss or metrics here.
model.compile(optimizer="adam")

# Just use `fit` as usual -- you can use callbacks, etc.
x = np.random.random((1000, 32))
y = np.random.random((1000, 1))
model.fit(x, y, epochs=5)
```

```
1.9270
Epoch 2/5
32/32 ━━━━━━━━━━━━━━━━ 0s 385us/step - loss: 2.2155 - mae:
1.3920
Epoch 3/5
32/32 ━━━━━━━━━━━━━━━━ 0s 336us/step - loss: 1.1863 - mae:
0.9700
Epoch 4/5
32/32 ━━━━━━━━━━━━━━━━ 0s 373us/step - loss: 0.6510 - mae:
0.6811
Epoch 5/5
32/32 ━━━━━━━━━━━━━━━━ 0s 330us/step - loss: 0.4059 - mae:
0.5094
<keras.src.callbacks.history.History at 0x2a7a02860>
```

Supporting `sample_weight` & `class_weight`

You may have noticed that our first basic example didn't make any mention of sample weighting. If you want to support the `fit()` arguments `sample_weight` and `class_weight`, you'd simply do the following:

- Unpack `sample_weight` from the `data` argument
- Pass it to `compute_loss` & `update_state` (of course, you could also just apply it manually if you don't rely on `compile()` for losses & metrics)
- That's it.



```
# Unpack the data. Its structure depends on your model and
# on what you pass to `fit()`.

if len(data) == 3:
    x, y, sample_weight = data
else:
    sample_weight = None
    x, y = data

with tf.GradientTape() as tape:
    y_pred = self(x, training=True) # Forward pass
    # Compute the loss value.
    # The loss function is configured in `compile()`.

    loss = self.compute_loss(
        y=y,
        y_pred=y_pred,
        sample_weight=sample_weight,
    )

    # Compute gradients
    trainable_vars = self.trainable_variables
    gradients = tape.gradient(loss, trainable_vars)

    # Update weights
    self.optimizer.apply(gradients, trainable_vars)

    # Update the metrics.
    # Metrics are configured in `compile()`.

    for metric in self.metrics:
        if metric.name == "loss":
            metric.update_state(loss)
        else:
            metric.update_state(y, y_pred, sample_weight=sample_weight)

    # Return a dict mapping metric names to current value.
    # Note that it will include the loss (tracked in self.metrics).
    return {m.name: m.result() for m in self.metrics}

# Construct and compile an instance of CustomModel
inputs = keras.Input(shape=(32,))
outputs = keras.layers.Dense(1)(inputs)
model = CustomModel(inputs, outputs)
model.compile(optimizer="adam", loss="mse", metrics=["mae"])

# You can now use sample_weight argument
x = np.random.random((1000, 32))
y = np.random.random((1000, 1))
sw = np.random.random((1000, 1))
model.fit(x, y, sample_weight=sw, epochs=3)
```

```
Epoch 1/3
32/32 ━━━━━━━━━━━━━━━━ 0s 2ms/step - mae: 0.4228 - loss:
0.1420
Epoch 2/3
32/32 ━━━━━━━━━━━━━━━━ 0s 449us/step - mae: 0.3751 - loss:
0.1058
Epoch 3/3
32/32 ━━━━━━━━━━━━━━━━ 0s 337us/step - mae: 0.3478 - loss:
0.0951

<keras.src.callbacks.history.History at 0x2a7491780>
```



in exactly the same way. Here's what it looks like:

```
class CustomModel(keras.Model):
    def test_step(self, data):
        # Unpack the data
        x, y = data
        # Compute predictions
        y_pred = self(x, training=False)
        # Updates the metrics tracking the loss
        loss = self.compute_loss(y=y, y_pred=y_pred)
        # Update the metrics.
        for metric in self.metrics:
            if metric.name == "loss":
                metric.update_state(loss)
            else:
                metric.update_state(y, y_pred)
        # Return a dict mapping metric names to current value.
        # Note that it will include the loss (tracked in self.metrics).
        return {m.name: m.result() for m in self.metrics}

# Construct an instance of CustomModel
inputs = keras.Input(shape=(32,))
outputs = keras.layers.Dense(1)(inputs)
model = CustomModel(inputs, outputs)
model.compile(loss="mse", metrics=["mae"])

# Evaluate with our custom test_step
x = np.random.random((1000, 32))
y = np.random.random((1000, 1))
model.evaluate(x, y)
```

```
32/32 ━━━━━━━━━━━━━━━━ 0s 927us/step - mae: 0.8518 - loss:
0.9166
[0.912325382232666, 0.8567370176315308]
```

Wrapping up: an end-to-end GAN example

Let's walk through an end-to-end example that leverages everything you just learned.

Let's consider:

- A generator network meant to generate 28x28x1 images.
- A discriminator network meant to classify 28x28x1 images into two classes ("fake" and "real").
- One optimizer for each.
- A loss function to train the discriminator.

```
    keras.Input(shape=(28, 28, 1)),
    layers.Conv2D(64, (3, 3), strides=(2, 2), padding="same"),
    layers.LeakyReLU(negative_slope=0.2),
    layers.Conv2D(128, (3, 3), strides=(2, 2), padding="same"),
    layers.LeakyReLU(negative_slope=0.2),
    layers.GlobalMaxPooling2D(),
    layers.Dense(1),
],
name="discriminator",
)

# Create the generator
latent_dim = 128
generator = keras.Sequential(
[
    keras.Input(shape=(latent_dim,)),
    # We want to generate 128 coefficients to reshape into a 7x7x128 map
    layers.Dense(7 * 7 * 128),
    layers.LeakyReLU(negative_slope=0.2),
    layers.Reshape((7, 7, 128)),
    layers.Conv2DTranspose(128, (4, 4), strides=(2, 2), padding="same"),
    layers.LeakyReLU(negative_slope=0.2),
    layers.Conv2DTranspose(128, (4, 4), strides=(2, 2), padding="same"),
    layers.LeakyReLU(negative_slope=0.2),
    layers.Conv2D(1, (7, 7), padding="same", activation="sigmoid"),
],
name="generator",
)
```

Here's a feature-complete GAN class, overriding `compile()` to use its own signature, and implementing the entire GAN algorithm in 17 lines in `train_step`:



```
super().__init__()  
self.discriminator = discriminator  
self.generator = generator  
self.latent_dim = latent_dim  
self.d_loss_tracker = keras.metrics.Mean(name="d_loss")  
self.g_loss_tracker = keras.metrics.Mean(name="g_loss")  
self.seed_generator = keras.random.SeedGenerator(1337)  
  
@property  
def metrics(self):  
    return [self.d_loss_tracker, self.g_loss_tracker]  
  
def compile(self, d_optimizer, g_optimizer, loss_fn):  
    super().compile()  
    self.d_optimizer = d_optimizer  
    self.g_optimizer = g_optimizer  
    self.loss_fn = loss_fn  
  
def train_step(self, real_images):  
    if isinstance(real_images, tuple):  
        real_images = real_images[0]  
    # Sample random points in the latent space  
    batch_size = tf.shape(real_images)[0]  
    random_latent_vectors = keras.random.normal(  
        shape=(batch_size, self.latent_dim), seed=self.seed_generator  
    )  
  
    # Decode them to fake images  
    generated_images = self.generator(random_latent_vectors)  
  
    # Combine them with real images  
    combined_images = tf.concat([generated_images, real_images], axis=0)  
  
    # Assemble labels discriminating real from fake images  
    labels = tf.concat(  
        [tf.ones((batch_size, 1)), tf.zeros((batch_size, 1))], axis=0  
    )  
    # Add random noise to the labels - important trick!  
    labels += 0.05 * keras.random.uniform(  
        tf.shape(labels), seed=self.seed_generator  
    )  
  
    # Train the discriminator  
    with tf.GradientTape() as tape:  
        predictions = self.discriminator(combined_images)  
        d_loss = self.loss_fn(labels, predictions)  
        grads = tape.gradient(d_loss, self.discriminator.trainable_weights)  
        self.d_optimizer.apply(grads, self.discriminator.trainable_weights)  
  
    # Sample random points in the latent space  
    random_latent_vectors = keras.random.normal(  
        shape=(batch_size, self.latent_dim), seed=self.seed_generator  
    )  
  
    # Assemble labels that say "all real images"  
    misleading_labels = tf.zeros((batch_size, 1))  
  
    # Train the generator (note that we should *not* update the weights  
    # of the discriminator)!  
    with tf.GradientTape() as tape:  
        predictions =  
    self.discriminator(self.generator(random_latent_vectors))  
        g_loss = self.loss_fn(misleading_labels, predictions)  
        grads = tape.gradient(g_loss, self.generator.trainable_weights)  
        self.g_optimizer.apply(grads, self.generator.trainable_weights)
```



```
self.g_loss_tracker.update_state(g_loss)
return {
    "d_loss": self.d_loss_tracker.result(),
    "g_loss": self.g_loss_tracker.result(),
}
```

Let's test-drive it:

```
# Prepare the dataset. We use both the training & test MNIST digits.
batch_size = 64
(x_train, _), (x_test, _) = keras.datasets.mnist.load_data()
all_digits = np.concatenate([x_train, x_test])
all_digits = all_digits.astype("float32") / 255.0
all_digits = np.reshape(all_digits, (-1, 28, 28, 1))
dataset = tf.data.Dataset.from_tensor_slices(all_digits)
dataset = dataset.shuffle(buffer_size=1024).batch(batch_size)

gan = GAN(discriminator=discriminator, generator=generator,
latent_dim=latent_dim)
gan.compile(
    d_optimizer=keras.optimizers.Adam(learning_rate=0.0003),
    g_optimizer=keras.optimizers.Adam(learning_rate=0.0003),
    loss_fn=keras.losses.BinaryCrossentropy(from_logits=True),
)

# To limit the execution time, we only train on 100 batches. You can train on
# the entire dataset. You will need about 20 epochs to get nice results.
gan.fit(dataset.take(100), epochs=1)
```

```
100/100 ━━━━━━━━━━━━━━━━ 51s 500ms/step - d_loss: 0.5645 -
g_loss: 0.7434
```

```
<keras.src.callbacks.history.History at 0x14a4f1b10>
```

The ideas behind deep learning are simple, so why should their implementation be painful?

[Terms](#) | [Privacy](#)